

JPA N+1

The JPA N+1 problem is a common issue encountered when using lazy loading with entity relationships in Java Persistence API (JPA). Here's a breakdown of what the N+1 problem is, why it occurs, and how to solve it, using a "Book" and "Author" example.

1. What is the N+1 Problem? [🔗](#)

- The N+1 problem arises when a main query (1) fetches a list of entities, and then additional (N) queries are executed for each entity to load related data.
- This often happens in one-to-many or many-to-many relationships where entities are loaded lazily.
- The main query fetches all records from the parent entity, and for each parent record, another query fetches the child entities, leading to **N+1** total queries.

2. Example Setup [🔗](#)

Let's consider two entities: **Book** and **Author**.

- A **Book** has one **Author**, but an **Author** can have multiple **Books**.

```
1 @Entity
2 public class Book {
3     @Id
4     private Long id;
5     private String title;
6
7     @ManyToOne(fetch = FetchType.LAZY)
8     @JoinColumn(name = "author_id")
9     private Author author;
10 }
11
12 @Entity
13 public class Author {
14     @Id
15     private Long id;
16     private String name;
17
18     @OneToMany(mappedBy = "author", fetch = FetchType.LAZY)
19     private List<Book> books;
20 }
```

3. How the N+1 Problem Occurs [🔗](#)

Consider fetching a list of books and their authors with the following query:

```
1 List<Book> books = entityManager.createQuery("SELECT b FROM Book b", Book.class).getResultList();
2
3 for (Book book : books) {
4     System.out.println(book.getAuthor().getName());
5 }
```

```
5 }
```

Here's what happens:

1. The first query fetches all `Book` records:

```
1 SELECT * FROM Book;
```

2. Then, for each `Book`, a separate query is executed to fetch its `Author`, resulting in `N` additional queries:

```
1 SELECT * FROM Author WHERE id = ?;
```

If there are 100 books, this results in $1 + 100 = 101$ queries, creating significant overhead and impacting performance.

4. Solutions to the N+1 Problem [🔗](#)

Solution 1: Fetch Join [🔗](#)

One way to avoid the N+1 problem is by using a `JOIN FETCH` to fetch associated entities in a single query.

```
1 List<Book> books = entityManager.createQuery(  
2     "SELECT b FROM Book b JOIN FETCH b.author", Book.class).getResultList();
```

In this query, `JOIN FETCH` ensures that `Author` entities are fetched along with each `Book`, eliminating the need for additional queries for each `Book`.

The resulting SQL will be a single query:

```
1 SELECT * FROM Book b INNER JOIN Author a ON b.author_id = a.id;
```

Solution 2: `@EntityGraph` Annotation [🔗](#)

Using the `@EntityGraph` annotation allows you to specify the relationships to fetch eagerly at query time without modifying the fetch type in the entity.

```
1 @EntityGraph(attributePaths = {"author"})  
2 @Query("SELECT b FROM Book b")  
3 List<Book> findAllBooks();
```

In this approach:

- Only one query is executed, which includes fetching both `Book` and `Author` entities.
- This is a flexible solution as it allows you to keep `LAZY` fetch type on the entity and specify eager loading only when needed.

Solution 3: Batch Fetching [🔗](#)

Batch fetching is another approach to optimize lazy loading by configuring JPA to load associated entities in batches rather than one-by-one.

- This can be configured in the `persistence.xml` file or with annotations.

```
1 <property name="hibernate.default_batch_fetch_size" value="10"/>
```

This way, if there are 100 `Books`, JPA will group the `Author` fetches in batches of 10, reducing the number of queries from 100 to 10.

5. Conclusion [↗](#)

The N+1 problem is a performance issue in JPA caused by lazy loading in entity relationships. To solve it:

1. Use `JOIN FETCH` to fetch associated entities in a single query.
2. Leverage `@EntityGraph` to load associations eagerly at query time.
3. Configure batch fetching to load data in chunks.

These approaches will reduce the number of queries, resulting in a more efficient application.